

# dl3

April 1, 2025

```
[1]: import numpy as np
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import RMSprop
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt
from sklearn import metrics
```

```
[2]: # Load the OCR dataset

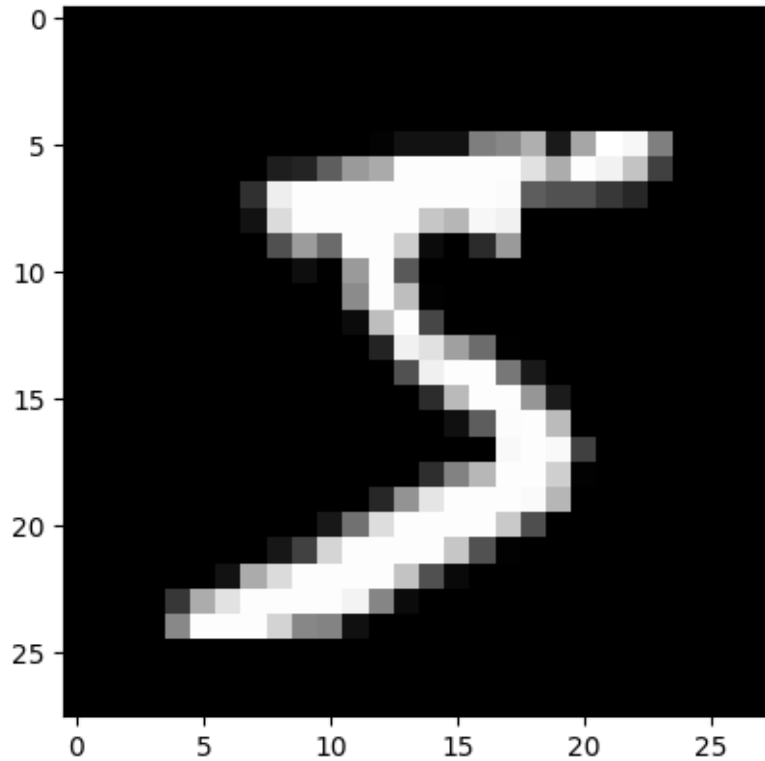
# The MNIST dataset is a built-in dataset provided by Keras.
# It consists of 70,000 28x28 grayscale images, each of which displays a single
    ↳ handwritten digit from 0 to 9.
# The training set consists of 60,000 images, while the test set has 10,000
    ↳ images.

(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434                      0s  
0us/step

```
[4]: # X_train and X_test are our array of images while y_train and y_test are our
    ↳ array of labels for each image.
# The first tuple contains the training set features (X_train) and the training
    ↳ set labels (y_train).
# The second tuple contains the testing set features (X_test) and the testing
    ↳ set labels (y_test).
# For example, if the image shows a handwritten 7, then the label will be the
    ↳ integer 7.

plt.imshow(x_train[0], cmap='gray') # imshow() function which simply displays
    ↳ an image.
plt.show() # cmap is responsible for mapping a specific colormap to the values
    ↳ found in the array that you passed as the first argument.
```



```
[5]: # image appears black and white and that each axis of the plot ranges from 0 to 28.
```

```
# This is because of the format that all the images in the dataset have:
```

```
# 1. All the images are grayscale, meaning they only contain black, white and grey.
```

```
# 2. The images are 28 pixels by 28 pixels in size (28x28).
```

```
print(x_train[0])
```

```
# image data is just an array of digits. You can almost make out a 5 from the pattern of the digits in the array.
```

```
# Array of 28 values
```

```
# a grayscale pixel is stored as a digit between 0 and 255 where 0 is black, 255 is white and values in between are different shades of gray.
```

```
# Therefore, each value in the [28][28] array tells the computer which color to put in that position when we display the actual image.
```

```
[ [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
    0  0  0  0  0  0  0  0  0  0  0 ]
  [ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
```

	0	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0]							
[	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0]							
[	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0	0]							
[	0	0	0	0	0	0	0	0	0	0	0	0	0	3	18	18	18	126
	175	26	166	255	247	127	0	0	0	0]								
[	0	0	0	0	0	0	0	0	30	36	94	154	170	253	253	253	253	253
	225	172	253	242	195	64	0	0	0	0]								
[	0	0	0	0	0	0	0	49	238	253	253	253	253	253	253	253	253	251
	93	82	82	56	39	0	0	0	0	0]								
[	0	0	0	0	0	0	0	18	219	253	253	253	253	253	198	182	247	241
	0	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	80	156	107	253	253	205	11	0	43	154
	0	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	14	1	154	253	90	0	0	0	0
	0	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	0	139	253	190	2	0	0	0
	0	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	0	11	190	253	70	0	0	0
	0	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	0	0	35	241	225	160	108	1
	0	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	0	0	0	81	240	253	253	119
	25	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	0	0	0	0	45	186	253	253
	150	27	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	16	93	252
	253	187	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	249
	253	249	64	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	0	0	0	0	46	130	183	253
	253	207	2	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	0	0	39	148	229	253	253	253
	250	182	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	24	114	221	253	253	253	253	201
	78	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	23	66	213	253	253	253	253	198	81	2	
	0	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	18	171	219	253	253	253	195	80	9	0	0	0
	0	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	55	172	226	253	253	253	253	244	133	11	0	0	0	0
	0	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	136	253	253	253	212	135	132	16	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0]								
[	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

```

    0  0  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0]
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]

```

```

[6]: print("X_train shape", x_train.shape)
      print("y_train shape", y_train.shape)
      print("X_test shape", x_test.shape)
      print("y_test shape", y_test.shape)

```

```

X_train shape (60000, 28, 28)
y_train shape (60000,)
X_test shape (10000, 28, 28)
y_test shape (10000,)

```

```

[7]: # X: Training data of shape (n_samples, n_features)
      # y: Training label values of shape (n_samples, n_labels)
      # 2D array of height and width, 28 pixels by 28 pixels will just become 784
      ↪ pixels (28 squared).
      # Remember that X_train has 60,000 elements, each with 784 total pixels so
      ↪ will become shape (60000, 784).
      # Whereas X_test has 10,000 elements, each with each with 784 total pixels so
      ↪ will become shape (10000, 784).

      x_train = x_train.reshape(60000, 784)
      x_test = x_test.reshape(10000, 784)
      x_train = x_train.astype('float32') # use 32-bit precision when training a
      ↪ neural network, so at one point the training data will have to be converted
      ↪ to 32 bit floats. Since the dataset fits easily in RAM, we might as well
      ↪ convert to float immediately.
      x_test = x_test.astype('float32')
      x_train /= 255 # Each image has Intensity from 0 to 255
      x_test /= 255

      # Regarding the division by 255, this is the maximum value of a byte (the input
      ↪ feature's type before the conversion to float32),
      # so this will ensure that the input features are scaled between 0.0 and 1.0.
      # USING svm-https://mgta.gmu.edu/courses/ml-with-python/
      ↪ handwrittenDigitRecognition.php#:~:
      ↪ text=Remember%20that%20X_train%20has%2060%2C0

```

```

[8]: # Convert class vectors to binary class matrices
      num_classes = 10
      y_train = np.eye(num_classes)[y_train] # Return a 2-D array with ones on the
      ↪ diagonal and zeros elsewhere.

```

```
y_test = np.eye(num_classes)[y_test] # f your particular categories is present
↳ then it mark as 1 else 0 in remain r
```

```
[9]: # Define the model architecture
model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(784,))) # The input_shape
↳ argument is passed to the foremost layer. It comprises of a tuple shape,
model.add(Dropout(0.2)) # DROP OUT RATIO 20%
model.add(Dense(512, activation='relu')) #returns a sequence of vectors of
↳ dimension 512
model.add(Dropout(0.2))
model.add(Dense(num_classes, activation='softmax'))
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87:  
UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When  
using Sequential models, prefer using an `Input(shape)` object as the first  
layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[10]: # Compile the model
model.compile(loss='categorical_crossentropy', # for a multi-class
↳ classification problem
              optimizer=RMSprop(),
              metrics=['accuracy'])
```

```
[11]: # Train the model
batch_size = 128 # batch_size argument is passed to the layer to define a batch
↳ size for the inputs.
epochs = 20
history = model.fit(x_train, y_train,
                    batch_size=batch_size,
                    epochs=epochs,
                    verbose=1, # verbose=1 will show you an animated progress
↳ bar eg. [=====]
                    validation_data=(x_test, y_test)) # Using validation_data
↳ means you are providing the training set and validation set yourself,
                                                    # validation_split means
↳ you only provide a training set and keras splits it into a training set and
↳ a validation set
```

Epoch 1/20

469/469 13s 25ms/step -

accuracy: 0.8615 - loss: 0.4496 - val\_accuracy: 0.9598 - val\_loss: 0.1299

Epoch 2/20

469/469 19s 21ms/step -

accuracy: 0.9664 - loss: 0.1092 - val\_accuracy: 0.9732 - val\_loss: 0.0896

Epoch 3/20

469/469 11s 23ms/step -  
accuracy: 0.9767 - loss: 0.0771 - val\_accuracy: 0.9775 - val\_loss: 0.0757  
Epoch 4/20

469/469 21s 24ms/step -  
accuracy: 0.9830 - loss: 0.0571 - val\_accuracy: 0.9791 - val\_loss: 0.0700  
Epoch 5/20

469/469 19s 22ms/step -  
accuracy: 0.9849 - loss: 0.0460 - val\_accuracy: 0.9838 - val\_loss: 0.0617  
Epoch 6/20

469/469 11s 23ms/step -  
accuracy: 0.9874 - loss: 0.0388 - val\_accuracy: 0.9835 - val\_loss: 0.0610  
Epoch 7/20

469/469 21s 24ms/step -  
accuracy: 0.9892 - loss: 0.0324 - val\_accuracy: 0.9812 - val\_loss: 0.0658  
Epoch 8/20

469/469 20s 23ms/step -  
accuracy: 0.9916 - loss: 0.0272 - val\_accuracy: 0.9824 - val\_loss: 0.0709  
Epoch 9/20

469/469 10s 21ms/step -  
accuracy: 0.9913 - loss: 0.0268 - val\_accuracy: 0.9823 - val\_loss: 0.0661  
Epoch 10/20

469/469 11s 23ms/step -  
accuracy: 0.9926 - loss: 0.0230 - val\_accuracy: 0.9831 - val\_loss: 0.0736  
Epoch 11/20

469/469 11s 23ms/step -  
accuracy: 0.9939 - loss: 0.0188 - val\_accuracy: 0.9840 - val\_loss: 0.0754  
Epoch 12/20

469/469 11s 24ms/step -  
accuracy: 0.9948 - loss: 0.0161 - val\_accuracy: 0.9823 - val\_loss: 0.0798  
Epoch 13/20

469/469 11s 23ms/step -  
accuracy: 0.9944 - loss: 0.0162 - val\_accuracy: 0.9841 - val\_loss: 0.0730  
Epoch 14/20

469/469 10s 22ms/step -  
accuracy: 0.9954 - loss: 0.0135 - val\_accuracy: 0.9849 - val\_loss: 0.0676  
Epoch 15/20

469/469 21s 23ms/step -  
accuracy: 0.9957 - loss: 0.0137 - val\_accuracy: 0.9819 - val\_loss: 0.0801  
Epoch 16/20

469/469 21s 24ms/step -  
accuracy: 0.9959 - loss: 0.0125 - val\_accuracy: 0.9848 - val\_loss: 0.0762  
Epoch 17/20

469/469 19s 21ms/step -  
accuracy: 0.9967 - loss: 0.0104 - val\_accuracy: 0.9850 - val\_loss: 0.0759  
Epoch 18/20

469/469 11s 24ms/step -  
accuracy: 0.9963 - loss: 0.0111 - val\_accuracy: 0.9847 - val\_loss: 0.0850  
Epoch 19/20

```
469/469          20s 23ms/step -  
accuracy: 0.9968 - loss: 0.0103 - val_accuracy: 0.9822 - val_loss: 0.0966  
Epoch 20/20  
469/469          11s 23ms/step -  
accuracy: 0.9968 - loss: 0.0104 - val_accuracy: 0.9842 - val_loss: 0.0831
```

```
[12]: # Evaluate the model  
score = model.evaluate(x_test, y_test, verbose=0)  
print('Test loss:', score[0])  
print('Test accuracy:', score[1])
```

```
Test loss: 0.08312281966209412  
Test accuracy: 0.9842000007629395
```

```
[ ]:
```