

In [3]: %%writefile add.cu

```
#include <iostream>
#include <cuda_runtime.h> // Provides necessary functions and macros to

using namespace std;

__global__ void addVectors(int* A, int* B, int* C, int n) //__global__: S
{
    int i = blockIdx.x * blockDim.x + threadIdx.x; // blockIdx.x: The ind
    if (i < n) // Ensures that threads do not access memory beyond the al
    {
        C[i] = A[i] + B[i]; // Adds corresponding elements from vectors A
    }
}

int main()
{
    int n = 1000000; // The size of the vectors (one million elements).
    int* A, * B, * C; // Pointers for the host (CPU) memory.
    int size = n * sizeof(int); // The memory size required for each vect

    // Allocate memory on the host
    cudaMallocHost(&A, size); // Allocates pinned (page-locked) memory o
    cudaMallocHost(&B, size);
    cudaMallocHost(&C, size);

    // Initialize the vectors
    for (int i = 0; i < n; i++)
    {
        A[i] = i; // Fills A with values [0, 1, 2, ..., n-1].
        B[i] = i * 2; // Fills B with values [0, 2, 4, ..., 2*(n-1)].
    }

    // Allocate memory on the device
    int* dev_A, * dev_B, * dev_C;
    cudaMalloc(&dev_A, size); // cudaMalloc(&dev_A, size): Allocates size
    cudaMalloc(&dev_B, size); // cudaMalloc(&dev_B, size): Allocates size
    cudaMalloc(&dev_C, size); // cudaMalloc(&dev_C, size): Allocates size

    // Copy data from host to device
    cudaMemcpy(dev_A, A, size, cudaMemcpyHostToDevice); // cudaMemcpy(des
    cudaMemcpy(dev_B, B, size, cudaMemcpyHostToDevice);

    // Launch the kernel
    int blockSize = 256; // Defines 256 threads per block.
    int numBlocks = (n + blockSize - 1) / blockSize; // Ensures that all
    addVectors<<<numBlocks, blockSize>>>(dev_A, dev_B, dev_C, n); // This

    // Copy data from device to host
    cudaMemcpy(C, dev_C, size, cudaMemcpyDeviceToHost); // Copies the com

    // Print the results
    for (int i = 0; i < 10; i++)
    {
        cout << C[i] << " "; // Prints the first 10 elements of C to veri
    }
    cout << endl;
```

```
// Free memory
cudaFree(dev_A); // releases memory on the GPU.
cudaFree(dev_B);
cudaFree(dev_C);
cudaFreeHost(A); // releases pinned memory on the CPU.
cudaFreeHost(B);
cudaFreeHost(C);

return 0;
}
```

Overwriting add.cu

```
In [4]: !rm -rf /usr/local/cuda          # Removes any previous CUDA installatio
!ln -s /usr/local/cuda-12.5 /usr/local/cuda    # Links to CUDA 12.2.
!nvcc -arch=sm_75 add.cu -o add                # Compiles the CUDA program (nvcc
```

```
In [5]: !./add // Each element of C is the sum of the corresponding elements of A
0 3 6 9 12 15 18 21 24 27
```

```
In [ ]:
```